

TIEU KIM HAO

haotieu1601@gmail.com • (+84) 333 983 245 • [linkedin.com/in/haotieu2001](https://www.linkedin.com/in/haotieu2001)
[haotieu2001.github.io](https://github.com/haotieu2001)

Data Engineer with 3+ years in data pipelines, cloud infrastructure and CI/CD. Now on a commercial bank's Databricks and AWS platform — daily ingestion from 19 sources into a 4.8 TB lakehouse, plus the reliability and cost optimization behind it.

WORK EXPERIENCE

Data Engineer

Jan 2026 - Present

Vietnam International Bank (VIB)

Ho Chi Minh City, Vietnam

Project: NEDP — Next-gen Enterprise Data Platform (Databricks Unity Catalog + Delta Lakehouse on AWS)

Responsibilities:

- Delivered 145 source tables (300 model files) from Oracle, PostgreSQL and SQL Server into a Delta lakehouse — 19 domains, roughly a quarter of a 1,200-model production estate.
- Federated source databases into Unity Catalog as foreign catalogs, so Spark queries them directly without a staging copy.
- Authored 21 Databricks Workflows: split one large daily pipeline into separate light and heavy jobs, and isolated unreliable sources so one failure no longer blocks the rest.
- Designed source-freshness gating for daily ingestion (fail / warn / non-blocking strategies plus upstream job-completion signals) so jobs wait on the real source SLA, not a fixed clock.
- Fixed a production Spark outage on EMR caused by reclaimed spot nodes: added executor exclusion so the driver stops sending work to dead executors, and moved shuffle data to S3 so it survives.
- Led a Delta Liquid Clustering migration on a 4.8 TB lakehouse, converting 90+ tables from a 108-table pilot and raising average file size 7×, from 5.4 to 37.2 MB.
- Designed an S3 cost program targeting a 50% storage cut on four production buckets, built in Terraform with lifecycle rules and Intelligent-Tiering.
- Cut EMR cost by capping on-demand fallback to the fleet baseline, and moved daily Databricks jobs onto right-sized job clusters.
- Built 6 Azure DevOps pipelines that validate and deploy Databricks jobs to dev, uat and prod.

Key Tools: Databricks (Unity Catalog, Delta, Asset Bundles), Spark, dbt, Python, SQL, Oracle, AWS (S3, EMR, Glue, IAM, MWAA), Terraform, Airflow, Azure DevOps

Data Engineer

Mar 2024 - Dec 2025

ConsensysAI (Singapore)

Remote

Project: Consensys World — a platform analyzing millions of consumer reviews for global market insights

Responsibilities:

- Built and ran a multi-process pipeline that collected ~1M review records a week from US, Brazil and Australia into AWS S3.
- Kept the pipeline running every day with retry, proxy rotation and CAPTCHA handling, even when sites slowed down or changed.
- Used Python and Pandas to clean every source into one format, then OpenAI APIs to label each review by sentiment and topic.
- Led a 4-member data collection team, assigning work, reviewing code and tracking weekly delivery to keep the dataset on schedule.

Key Tools: Python, Pandas, AWS S3, Playwright, Requests, BeautifulSoup, OpenAI APIs, Git, Linux CLI

DevOps Engineer

Jul 2023 - Feb 2024

OPSWAT

Ho Chi Minh City, Vietnam

Project: MetaDefender Core — enterprise malware detection and file sanitization engine

Responsibilities:

- Built and maintained automated CI pipelines in Jenkins and TeamCity, streamlining code integration and deployment.
- Replicated customer environments on AWS and ran a local Kubernetes cluster for reproducible testing and development.
- Maintained and enhanced Python-based performance testing tools, monitoring throughput and stability for critical processes.

Key Tools: AWS, Terraform, Jenkins, TeamCity, Docker, Kubernetes, Python, Linux, Git

EDUCATION

Bachelor of DATA SCIENCE

Sep 2019 - Jun 2023

University of Information Technology - VNUHCM

GPA: 3.3/4.0

SKILLS

Data Engineering: Databricks (Unity Catalog, Delta Lake), Spark/PySpark, dbt, SQL, Apache Airflow

Data Quality & Governance: source-freshness gating, reconciliation, pipeline monitoring, Unity Catalog grants, cost tagging

Cloud & Infrastructure: AWS (S3, EMR, Glue, IAM, MWAA, EC2, Lambda), Terraform, Docker, Kubernetes

Programming & Tooling: Python, Bash scripting, Git, Azure DevOps / Jenkins CI-CD

CERTIFICATIONS

Databricks Certified Machine Learning Professional

Jul 2026 - Jul 2028

IELTS 6.5

Nov 2022

LANGUAGES

English (Professional working proficiency)